

Elastic Compute Cloud- EC2

EC2- Getting Started

- **Elastic Compute Cloud**
- Provides **resizable compute** capacity, **Pay** per use
- One of the most used AWS service and core service
- Enables you to build and run applications faster
- Can scale up or down when there is an increase or decrease in website traffic



EC2- Provision Options

- **On demand** – For Short workload, pay fixed rate per hour (per second), no upfront cost, In customer control
- **Reserved Instances** – Reserved capacity in AWS datacenter, duration 1 to 3 years, and flexible payment options, more discount
 - **Standard RI** – Can not change instance attributes like types, OS in between, more discount as compared to Convertible RI
 - **Convertible RI** – Change Instance types, OS etc in between, flexibility in instance types, expensive than standard RI
 - **Scheduled Instances** -Instances that are always available on the specified recurring schedule, for a one-year term e.g. Every Friday 5 PM – 8 PM. Not available in all region
- **Spot** – Bid instances, if application has flexible start and end time, more discount up to 90 % compared to On-Demand, Spot pricing based on offer and capacity availability at AWS
 - <https://aws.amazon.com/ec2/spot/pricing/>
- **Dedicated hosts**
- **Dedicated Instances**

EC2 – Instance Types

- Comprise varying combinations of CPU, memory, storage, and networking capacity
- Provides flexibility to choose the appropriate mix of resources for your applications

[PAGE CONTENT](#)

[General Purpose](#)

[Compute Optimized](#)

[Memory Optimized](#)

[Accelerated Computing](#)

[Storage Optimized](#)

[HPC Optimized](#)

[Instance Features](#)

[Measuring Instance Performance](#)

Instance Size	vCPU	Memory (GiB)	Instance Storage (GB)	Network Bandwidth (Gbps)**	EBS Bandwidth (Gbps)
c7g.medium	1	2	EBS-Only	Up to 12.5	Up to 10
c7g.large	2	4	EBS-Only	Up to 12.5	Up to 10
c7g.xlarge	4	8	EBS-Only	Up to 12.5	Up to 10
c7g.2xlarge	8	16	EBS-Only	Up to 15	Up to 10
c7g.4xlarge	16	32	EBS-Only	Up to 15	Up to 10
c7g.8xlarge	32	64	EBS-Only	15	10

EC2- AMI

- AMI (Amazon Machine Image)
 - **Software** and **hardware** configuration when instance launched
 - **Sources**
 - New Image from scratch
 - Published by AWS
 - AWS marketplace
 - Generated from existing instance
 - Uploaded Virtual Servers
- **Owned by me** – Specifically owned by your AWS account
- **Public Images** – Shared with all AWS accounts
- **Private images** – Shared with your AWS account and your private ones

EC2- Security Groups

- **Security Groups**
 - **Virtual** firewall
 - Apply on **instance** level
 - **Stateful** – No need to add the rules explicitly in outbound
 - 1 Instance can have many SGs
 - One SG can have many Rules
 - Rules get applied **immediately**
 - Can not add rules to **deny** or **block** specific IP address
 - Its **whitelist** - Anything not listed there will be blocked

EC2 (Security Groups)- LAB

EC2 Lab- Provision Linux VM

- Provision Linux instance with below configuration and Details
 - With additional 5 GB encrypted disk
 - Install OS update patches
 - Create new Security Group which allows SSH and http
 - Install Apache (httpd)
 - Launch simple html web page running on port 80
 - Stop instance
 - Terminate instance

EC2 Lab- Provision Windows VM

- Provision Windows instance with below configuration and Details
 - Run any bootstrap code
 - Add additional Disk
 - Create new Security Group which allows RDP from your IP only
 - Stop instance
 - Terminate instance

EC2 (AMI)- LAB

EC2- Snapshots

- **Snapshots**
 - Kind of backup/recovery
 - Protect your data
 - Data for snapshots stored on S3
 - Need snapshot restore on different region?- then copy the snapshot and create volume
 - Can increase the size of volume during creation
 - Snapshot Life Cycle Policy Or Data Lifecycle Manager

EC2 (Snapshot)- LAB

AWS User Data / Custom Scripts

- Can provide shell script or cloud-directives code
- Option of passing user data to the instance that can be used to perform common **automated configuration tasks**
- Run scripts after the instance **first starts (boots up)**
- Configure instance and install applications **programmatically**
- Also, you can use **AWS CLI** to run the commands
- User data scripts and cloud-init directives run only during the first boot cycle when an instance is launched
- Can configure script to be run at **instance restart**
- On **Windows**, for every restart
 - `<persist>true</persist>`
- On **Linux**, for every restart
 - `#cloud-config cloud_final_modules: - [scripts-user, always]`

User Data – Shell Script Example

```
#!/bin/bash
yum update -y
yum install httpd -y
cd /var/www/html
echo "My web server"> index.html
service httpd start
chkconfig httpd on
```

EC2 User Data - LAB


EC2 (Load Balancer)

EC2 Load Balancer – In general

- Virtual Entity or **appliance**
- To **balance** the load or traffic across EC2 instances
- Load balancer serves as the single point of contact for clients
- Distributes incoming application traffic across multiple EC2 instances in multiple Availability Zones
- Enable to achieve **high availability** for application
- Also, scale in and scale out automatically as per demand if you use this with Autoscaling
- Supports HTTP, HTTPS, TCP and SSL protocols
- **Types**
 - **Application**
 - **Network**
 - **Gateway LB (New Launch)**
- <https://aws.amazon.com/elasticloadbalancing/features/#compare>

Application Load Balancer

- Routing requests to multiple applications on a single EC2 instance
- Can configure specific request goes for specific content to specific instance
- <https://docs.aws.amazon.com/elasticloadbalancing/latest/application/introduction.html>
- Support multiple listeners with multiple target groups
- Functions at the application layer.
i.e. 7th layer of OSI model
- Improved performance
- Only support for VPC platform



EC2 (Load Balancer)- LAB


Labs

- Provision Win instance with bootstrap code
- Provision Linux instance with bootstrap code
- Elastic IP, Public IP, Private IP Demo
- CLI
 - How to set CLI on your computer
 - How to run CLI directly from AWS instance
 - EC2 operations- start, stop, terminate etc
 - S3- copy from one region to other, download files
 - EC2 Role
- Provision Linux machine. Install Apache. Copy data from S3 and launch webpage without Role.
- Provision Linux with custom scripts with Role


EC2-Autoscaling

- Can also be called as Auto-healing
- Helps you ensure that you have the correct number of Amazon EC2 instances available to handle the load for your application
- Auto Scaling monitors your applications and automatically adjusts capacity to maintain steady, predictable performance at the lowest possible cost
- Specify min and max number of instances in group
- No additional fees with Amazon EC2 Auto Scaling configuration

EC2- Autoscaling



Auto-scaling



EC2 AutoScaling - LAB

IP – Private/Public/Elastic

- Private
 - Based on vpc/subnet in which EC2 is setup
 - Upon restart, remains same
 - Like your nickname
- Public
 - Allocated from a pool of available public IP's
 - Upon restart, may get change
 - Like your mobile number
- Elastic
 - Static public IP
 - Upon restart, remains same
 - Like your unique Adhaar number

EC2 - Pricing

- Depend on which provisioning model
 - Saving plans, On Demand, RI, Scheduled Instances, Dedicated Instance/Dedicated host, Marketplace, Spot Instance etc
- Pay only for what you use
- Costing Attributes – Running state, storage, data transfer
- Display price an hourly rate – which instance types you are using, hourly or seconds
- Across AZ/region data transfer
- For stop instance – Only volume or storage (EBS) cost, no running and data transfer cost
- AWS Free Tier - 750 hours of Linux and Windows t2.micro instances, (t3.micro for the regions in which t2.micro is unavailable) each month for one year
- <https://aws.amazon.com/ec2/pricing/>